

# Lab 10 - Concurrency 2

*Out: December 3-4, 2012*

## 1 Introduction

Running programs with multiple threads can serve two different purposes:

- *performance*: Using multiple threads can increase performance in a program by allowing other threads to run while one thread sleeps, which is often the case with programs that have to access the network or disk. On a computer with multiple procesors, using multiple threads can help take advantage of that.
- *abstraction*: Using multiple threads in a program can also serve as a convenient abstraction. For example, a server that has to deal with multiple clients may want to have a different thread for each client. This way, each thread can operate independently, dealing with only one client at a time.

The issue with multithreaded programs, though, is that any data structures that are shared between the threads must be modified carefully, since each thread could be modifying the same element at the same time. As discussed in class, mutual exclusion is one way of solving this. In this assignment, we have provided a simple implementation of a linked list, along with a program that launches a number of threads that modify this list concurrently. The linked list implementation currently has no measures to provide thread-safety. Your task is to modify this code in order to make it safe for multithreading.

Two general approaches to making data structures thread-safe with mutual exclusion are *coarse-grained locking* and *fine-grained locking*. Coarse-grained locking involves locking the entire data structure with one lock, so that only one thread can access the data structure at any time. Fine-grained locking involves locking each component of the data structure separately. This approach allows multiple threads to access the data structure at the same time. In this lab, you will be implementing fine-grained locking, since implementing coarse-grained locking is fairly trivial in this case.

## 2 Assignment

Your assignment is to modify the code in *linkedlist.c* to be thread safe.

The functions in this file are:

- **main()**: Takes care of reading command line input and launching threads.
- **randomListManip()**: Randomly inserts or removes values from the linked list. The method's signature reflects the fact that it is being used in `pthread_create` to launch a new thread. The given *Makefile* generates a binary called *random* that uses this method.
- **seqListManip()**: Inserts values into the linked list. Running this method in multiple threads concurrently makes the effects of unsafe multithreaded programming quite apparent. Since

each thread only adds elements to the list and never removes elements, anytime this method is run concurrently we can calculate the number of elements that should be in the list at the end of the program. The method's signature reflects the fact that it is being used in `pthread_create` to launch a new thread. The given *Makefile* generates a binary called *sequential* that uses this method.

- `search()`: Searches through the linked list for a node with the given value.
- `insertList()`: Inserts an element into the linked list if it is not already in the linked list.
- `deleteList()`: Deletes an element from the linked list if it exists in the linked list.

Running `make` using the provided *Makefile* creates two binaries: *sequential* and *random*. The former binary gives the `seqListManip` function as an argument to `pthread_create`, and the latter does the same with `randomListManip` instead. However, there is only one source file you will have to modify: *linkedlist.c*. The different binaries are generated by the macro surrounding the line that makes calls to `pthread_create`, which simply swaps which function is being given to the new threads to run.

## 2.1 Implementation Details

If you were implementing coarse-grained locking, it would be enough to modify the `seqListManip()` or `randomListManip()` functions. However, since you are implementing fine-grained locking, you will have to modify the `search()`, `insertList()`, and `deleteList()` functions. You will want to lock each node in the linked list separately, so that only one thread at a time can access the fields of any node in the linked list.

## 2.2 Tips: Using pthreads

You will probably want to use or learn more about the following pthread library functions in your implementation. Use the `man` command to find out more about them.

- `pthread`
  - `pthread_create`
  - `pthread_join`
  - `pthread_exit`
- `pthread_mutex`
  - `pthread_mutex_init`
  - `pthread_mutex_destroy`
  - `pthread_mutex_lock`
  - `pthread_mutex_unlock`
- `pthread_barrier`
  - `pthread_barrier_init`
  - `pthread_barrier_wait`

In addition to that, it may be useful to look up the use of `PTHREAD_MUTEX_INITIALIZER` to statically initialize a mutex.

## 2.3 Testing

Start testing your program by using the *sequential* binary. This is a simple case where each thread adds a sequence of numbers to the list, though the sequences for each thread overlap. When running this on the stencil code (which is not thread-safe), the resulting linked list varies in size each time the program is run. You can get a count of the number of elements in the list at the end by running `./sequential <nthreads> <niterations> 2>/dev/null | wc -l`. Since this implementation is unsafe, some duplicates arise in the list, leading to lists that are larger than they should be in a thread-safe implementation.

As soon as you are sure that your program works with this simpler binary, move on to the *random* binary.

The program in the *random* binary adds and removes elements at random from the list, and will be a better test of whether your locking is valid. Do note that since the *sequential* binary is deterministic, it will not test your code fully, and you should move on to using the *random* binary before getting checked off.

## 2.4 Debugging

Concurrent programming occasionally leads to a common class of bugs called deadlocks. To solve deadlocks, it is often helpful to look at the state of threads while they are stuck to identify what they are waiting for. `gdb` offers a set of helpful debugging tools for multithreaded programs that allows just that. To start `gdb` with your program, type

```
gdb ./random
```

Once `gdb` is started up, use the `run` command followed by the program arguments to run your code.

While your program is running, it is possible to interrupt and pause it so you may inspect the state of each thread. You may interrupt `gdb` using `CTRL-Z`.

From here, you may print a list of threads by typing `info threads`. This command enumerates the threads. You will see a `*` next to the current thread you are inspecting. To inspect a different thread, type `thread NUMBER` where `NUMBER` corresponds to the numbers in the list you printed.

You will find it is most helpful to print a backtrace with the `bt` command. This will print the functions the current thread is executing.

Use this information to discern where any deadlocks are coming from and how to fix your code.

## 3 Getting Checked Off

Once you've completed the lab, call a TA over and have them inspect your work. If you do not complete the lab during your lab session, you can get credit for it by completing it on your own

time and bringing it to TA hours before next weeks lab. Remember to read the course missive for information about course requirements and policies regarding labs and assignments.