

CS 33

Machine Programming (5)

All of the slides in this lecture are either from or adapted from slides provided by the authors of the textbook “Computer Systems: A Programmer’s Perspective,” 2nd Edition and are provided from the website of Carnegie-Mellon University, course 15-213, taught by Randy Bryant and David O’Hallaron in Fall 2010. These slides are indicated “Supplied by CMU” in the notes section of the slides.

Today

- **Arrays**
 - one-dimensional
 - multi-dimensional (nested)
 - multi-level
- **Structures**
 - allocation
 - access
 - alignment

Supplied by CMU.

Basic Data Types

- **Integral**

- stored & operated on in general (integer) registers
- signed vs. unsigned depends on instructions used

Intel	ASM	Bytes	C
byte	b	1	[unsigned] char
word	w	2	[unsigned] short
double word	l	4	[unsigned] int
quad word	q	8	[unsigned] long int (x86-64)

- **Floating Point**

- stored & operated on in floating point registers

Intel	ASM	Bytes	C
single	s	4	float
double	l	8	double
extended	t	10/12/16	long double


Supplied by CMU.

Array Allocation

- Basic Principle

$T\ A[L];$

- array of data type T and length L
- contiguously allocated region of $L * \text{sizeof}(T)$ bytes




Supplied by CMU.

Array Access

- Basic Principle

$T\ A[L];$

- array of data type T and length L
- identifier A can be used as a pointer to array element 0: type T^*




Reference	Type	Value
<code>val[4]</code>	<code>int</code>	3
<code>val</code>	<code>int *</code>	x
<code>val+1</code>	<code>int *</code>	$x+4$
<code>&val[2]</code>	<code>int *</code>	$x+8$
<code>val[5]</code>	<code>int</code>	??
<code>*(val+1)</code>	<code>int</code>	5
<code>val + i</code>	<code>int *</code>	$x+4i$

Supplied by CMU.

Array Example

```
#define ZLEN 5
typedef int zip_dig[ZLEN];

zip_dig brown = { 1, 5, 2, 1, 3 };
zip_dig yale = { 0, 2, 1, 3, 9 };
zip_dig penn = { 9, 4, 7, 2, 0 };
```



- Declaration “zip_dig brown” equivalent to “int brown[5]”
- Example arrays were allocated in successive 20-byte blocks
 - not guaranteed to happen in general

Supplied by CMU.

Array Accessing Example

zip_dig brown;

1	5	2	1	3
---	---	---	---	---

16 20 24 28 32 36

```
int get_digit
(zip_dig z, int dig)
{
    return z[dig];
}
```

IA32

```
# %edx = z
# %eax = dig
movl (%edx,%eax,4),%eax # z[dig]
```

- Register `%edx` contains starting address of array
- Register `%eax` contains array index
- Desired digit at $4 * \%eax + \%edx$
- Use memory reference $(\%edx, \%eax, 4)$

Supplied by CMU.

Array Loop Example (IA32)


```
void zincr(zip_dig z) {  
    int i;  
    for (i = 0; i < ZLEN; i++)  
        z[i]++;  
}
```

```
# edx = z  
movl $0, %eax          # %eax = i  
.L4:                   # loop:  
    addl $1, (%edx,%eax,4) # z[i]++  
    addl $1, %eax         # i++  
    cmpl $5, %eax         # i:5  
    jne .L4              # if !=, goto loop
```

Supplied by CMU.

Nested-Array Example

```
#define PCOUNT 4
zip_dig pvd[PCOUNT] =
    {{1, 5, 2, 0, 6},
     {1, 5, 2, 1, 3 },
     {1, 5, 2, 1, 7 },
     {1, 5, 2, 2, 1 }};
```



- “zip_dig pvd[4]” equivalent to “int pvd[4][5]”
 - variable pvd: array of 4 elements, allocated contiguously
 - each element is an array of 5 int’s, allocated contiguously
- “Row-Major” ordering of all elements guaranteed


Supplied by CMU.

Multidimensional (Nested) Arrays

- **Declaration**
 - T $A[R][C]$;
 - 2D array of data type T
 - R rows, C columns
 - type T element requires K bytes
- **Array Size**
 - $R * C * K$ bytes
- **Arrangement**
 - row-major ordering

$$\begin{bmatrix} A[0][0] & \dots & A[0][C-1] \\ \vdots & & \vdots \\ A[R-1][0] & \dots & A[R-1][C-1] \end{bmatrix}$$

int $A[R][C]$;




Supplied by CMU.

Nested-Array Row Access

- Row Vectors

- $A[i]$ is array of C elements
- each element of type T requires K bytes
- starting address $A + i * (C * K)$

```
int A[R][C];
```



Supplied by CMU.

Nested-Array Row-Access Code

```
int *get_pvd_zip(int index)
{
    return pvd[index];
}
```

```
#define PCOUNT 4
zip_dig pvd[PCOUNT] =
    {{1, 5, 2, 0, 6},
     {1, 5, 2, 1, 3 },
     {1, 5, 2, 1, 7 },
     {1, 5, 2, 2, 1 }};
```

```
# %eax = index
leal (%eax,%eax,4),%eax # 5 * index
leal pvd(,%eax,4),%eax # pvd + (20 * index)
```

- **Row Vector**
 - pvd[index] is array of 5 int's
 - starting address pvd+20*index
- **IA32 Code**
 - computes and returns address
 - compute as pvd + 4*(index+4*index)


Supplied by CMU.

Nested-Array Element Access

- Array Elements

- $A[i][j]$ is element of type T , which requires K bytes
- address $A + i * (C * K) + j * K = A + (i * C + j) * K$

```
int A[R][C];
```



Supplied by CMU.

Nested-Array Element-Access Code

```
int get_pvd_digit
(int index, int dig)
{
    return pvd[index][dig];
}
```

```
movl    8(%ebp), %eax    # index
leal    (%eax,%eax,4), %eax # 5*index
addl    12(%ebp), %eax    # 5*index+dig
movl    pvd(,%eax,4), %eax # offset 4*(5*index+dig)
```

- **Array Elements**
 - pvd[index][dig] is int
 - address: $\text{pvd} + 20 \cdot \text{index} + 4 \cdot \text{dig}$
= $\text{pvd} + 4 \cdot (5 \cdot \text{index} + \text{dig})$
- **IA32 Code**
 - computes address $\text{pvd} + 4 \cdot ((\text{index} + 4 \cdot \text{index}) + \text{dig})$


Supplied by CMU.

Multi-Level Array Example

```
zip_dig brown = { 1, 5, 2, 1, 3 };  
zip_dig yale = { 0, 2, 1, 3, 9 };  
zip_dig penn = { 9, 4, 7, 2, 0 };
```

```
#define UCOUNT 3  
int *univ[UCOUNT] = {brown, yale, penn};
```

- Variable `univ` denotes array of 3 elements
- Each element is a pointer – 4 bytes
- Each pointer points to array of `int`'s



Supplied by CMU.

Element Access in Multi-Level Array

```
int get_univ_digit
(int index, int dig)
{
    return univ[index][dig];
}
```

```
movl 8(%ebp), %eax      # index
movl univ(,%eax,4), %edx # p = univ[index]
movl 12(%ebp), %eax     # dig
movl (%edx,%eax,4), %eax # p[dig]
```


- **Computation (IA32)**
 - element access $\text{Mem}[\text{Mem}[\text{univ} + 4 * \text{index}] + 4 * \text{dig}]$
 - must do two memory reads
 - » first get pointer to row array
 - » then access element within array

Supplied by CMU.

Array-Element Accesses


Nested array

```
int get_pvd_digit
(int index, int dig)
{
    return pvd[index][dig];
}
```



Multi-level array

```
int get_univ_digit
(int index, int dig)
{
    return univ[index][dig];
}
```



Accesses look similar in C, but addresses very different:

`Mem[pvd+20*index+4*dig]`

`Mem[Mem[univ+4*index]+4*dig]`

Supplied by CMU.

N X N Matrix Code

- **Fixed dimensions**

- know value of N at compile time

```
#define N 16
typedef int fix_matrix[N][N];
/* Get element a[i][j] */
int fix_ele
(fix_matrix a, int i, int j)
{
    return a[i][j];
}
```

- **Variable dimensions**

- just recently supported by gcc

```
/* Get element a[i][j] */
int var_ele
(int n, int a[n][n], int i, int j) {
    return a[i][j];
}
```

16 X 16 Matrix Access

- **Array Elements**

- address $A + i * (C * K) + j * K$
- $C = 16, K = 4$

```
/* Get element a[i][j] */  
int fix_ele(fix_matrix a, int i, int j) {  
    return a[i][j];  
}
```

```
movl 12(%ebp), %edx    # i  
sall $6, %edx          # i*64  
movl 16(%ebp), %eax    # j  
sall $2, %eax          # j*4  
addl 8(%ebp), %eax     # a + j*4  
movl (%eax,%edx), %eax # *(a + j*4 + i*64)
```

Supplied by CMU.

n X n Matrix Access

- **Array Elements**


- address $A + i * (C * K) + j * K$
- $C = n, K = 4$

```
/* Get element a[i][j] */  
int var_ele(int n, int a[n][n], int i, int j) {  
    return a[i][j];  
}
```

```
movl    8(%ebp), %eax    # n  
sall    $2, %eax        # n*4  
movl    %eax, %edx      # n*4  
imull   16(%ebp), %edx   # i*n*4  
movl    20(%ebp), %eax   # j  
sall    $2, %eax        # j*4  
addl    12(%ebp), %eax   # a + j*4  
movl    (%eax,%edx), %eax # *(a + j*4 + i*n*4)
```

Supplied by CMU.

Optimizing Fixed-Array Access



- **Computation**
 - step through all elements in column j
- **Optimization**
 - retrieving successive elements from single column

```
#define N 16
typedef int fix_matrix[N][N];

/* Retrieve column j from array */
void fix_column
(fix_matrix a, int j, int *dest)
{
    int i;
    for (i = 0; i < N; i++)
        dest[i] = a[i][j];
}
```

Optimizing Fixed-Array Access

- **Optimization**

- compute `aijp = &a[i][j]`
 - » initially = `a + 4*j`
 - » increment by `4*N`

Register	Value
<code>%ecx</code>	<code>aijp</code>
<code>%ebx</code>	<code>dest</code>
<code>%edx</code>	<code>i</code>

```
/* Retrieve column j from array */  
void fix_column  
(fix_matrix a, int j, int *dest)  
{  
    int i;  
    for (i = 0; i < N; i++)  
        dest[i] = a[i][j];  
}
```

```
.L8:                                # loop:  
    movl (%ecx), %eax              # Read *aijp  
    movl %eax, (%ebx,%edx,4)       # Save in dest[i]  
    addl $1, %edx                  # i++  
    addl $64, %ecx                 # aijp += 4*N  
    cmpl $16, %edx                 # i:N  
    jne .L8                        # if !=, goto loop
```

Supplied by CMU.

Optimizing Variable-Array Access

– Compute $ajp = \&a[i][j]$

- » initially = $a + 4*j$
- » increment by $4*n$

Register	Value
%ecx	aijp
%edi	dest
%edx	i
%ebx	$4*n$
%esi	n

```
/* Retrieve column j from array */
void var_column
(int n, int a[n][n],
 int j, int *dest)
{
    int i;
    for (i = 0; i < n; i++)
        dest[i] = a[i][j];
}
```

```
.L18:                                # loop:
    movl (%ecx), %eax               # Read *aijp
    movl %eax, (%edi,%edx,4)         # Save in dest[i]
    addl $1, %edx                   # i++
    addl $ebx, %ecx                 # aijp += 4*n
    cmpl $edx, %esi                 # n:i
    jg .L18                        # if >, goto loop
```

Supplied by CMU.

Today


- **Arrays**
 - one-dimensional
 - multi-dimensional (nested)
 - multi-level
- **Structures**
 - **allocation**
 - **access**
 - **alignment**

Supplied by CMU.

Structure Allocation

```
struct rec {  
    int a[3];  
    int i;  
    struct rec *n;  
};
```


Memory layout



- **Concept**
 - contiguously allocated region of memory
 - refer to members within structure by names
 - members may be of different types

Structure Access

```
struct rec {  
    int a[3];  
    int i;  
    struct rec *n;  
};
```



- Accessing structure member
 - pointer indicates first byte of structure
 - access elements with offsets


```
void  
set_i(struct rec *r,  
      int val)  
{  
    r->i = val;  
}
```

IA32 Assembly

```
# %edx = val  
# %eax = r  
movl %edx, 12(%eax) # Mem[r+12] = val
```

Generating Pointer to Structure Member

```
struct rec {  
    int a[3];  
    int i;  
    struct rec *n;  
};
```



- **Generating pointer to array element**
 - offset of each structure member determined at compile time
 - arguments
 - » `Mem[%ebp+8]: r`
 - » `Mem[%ebp+12]: idx`

```
int *get_ap  
(struct rec *r, int idx)  
{  
    return &r->a[idx];  
}
```


```
movl    12(%ebp), %eax    # Get idx  
sall    $2, %eax         # idx*4  
addl    8(%ebp), %eax     # r+idx*4
```

Following Linked List

- C code

```
void set_val
(struct rec *r, int val)
{
    while (r) {
        int i = r->i;
        r->a[i] = val;
        r = r->n;
    }
}
```

```
struct rec {
    int a[3];
    int i;
    struct rec *n;
};
```



Register	Value
%edx	r
%ecx	val

```
.L17:                                # loop:
    movl    12(%edx), %eax            # r->i
    movl    %ecx, (%edx,%eax,4)       # r->a[i] = val
    movl    16(%edx), %edx           # r = r->n
    testl   %edx, %edx               # Test r
    jne     .L17                     # If != 0 goto loop
```

Supplied by CMU.


Today

- **Arrays**
 - one-dimensional
 - multi-dimensional (nested)
 - multi-level
- **Structures**
 - allocation
 - access
 - **alignment**

Supplied by CMU.

Structures & Alignment


- Unaligned data



```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

- Aligned data

- primitive data type requires K bytes
- address must be multiple of K



Supplied by CMU.

Note that in this and the following slides, $p+i$ should be interpreted as meaning not $p[i]$, as it would be in C, but the address p plus the integer value i . Thus if p is, for example, $0x1004$, $p+1$ is $0x1005$.

Alignment Principles

- **Aligned data**
 - primitive data type requires K bytes
 - address must be multiple of K
 - required on some machines; advised on IA32
 - » treated differently by IA32 Linux, x86-64 Linux, and Windows!
- **Motivation for aligning data**
 - memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
 - » inefficient to load or store datum that spans quad word boundaries
 - » virtual memory very tricky when datum spans 2 pages
- **Compiler**
 - inserts gaps in structure to ensure correct alignment of fields

Supplied by CMU.

Specific Cases of Alignment (IA32)

- **1 byte:** `char`, ...
 - no restrictions on address
- **2 bytes:** `short`, ...
 - lowest 1 bit of address must be 0_2
- **4 bytes:** `int`, `float`, `char *`, ...
 - lowest 2 bits of address must be 00_2
- **8 bytes:** `double`, ...
 - Windows (and most other OS's & instruction sets):
 - » lowest 3 bits of address must be 000_2
 - Linux:
 - » lowest 2 bits of address must be 00_2
 - » i.e., treated the same as a 4-byte primitive data type
- **12 bytes:** `long double`
 - Windows, Linux:
 - » lowest 2 bits of address must be 00_2
 - » i.e., treated the same as a 4-byte primitive data type

Supplied by CMU.

Specific Cases of Alignment (x86-64)


- **1 byte:** `char`, ...
 - no restrictions on address
- **2 bytes:** `short`, ...
 - lowest 1 bit of address must be 0_2
- **4 bytes:** `int`, `float`, ...
 - lowest 2 bits of address must be 00_2
- **8 bytes:** `double`, `char *`, ...
 - Windows & Linux:
 - » lowest 3 bits of address must be 000_2
- **16 bytes:** `long double`
 - Linux:
 - » lowest 3 bits of address must be 000_2
 - » i.e., treated the same as a 8-byte primitive data type

Supplied by CMU.

Alignment with Structures

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

- **Within structure:**
 - must satisfy each element's alignment requirement
- **Overall structure placement**
 - each structure has alignment requirement K
 - » K = largest alignment of any element
 - initial address & structure length must be multiples of K
- **Example (under Windows or x86-64):**
 - $K = 8$, due to `double` element




Supplied by CMU.


Different Alignment Conventions

- **x86-64 or IA32 Windows:**
 - $K = 8$, due to double element

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```



- **IA32 Linux**
 - $K = 4$; double treated like a 4-byte data type




Supplied by CMU.

Meeting Overall Alignment Requirement

- For largest alignment requirement K
 - overall length of structure must be multiple of K

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```




Supplied by CMU.

Arrays of Structures

- Overall structure length is multiple of K
 - satisfy alignment requirement for every element

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```




Supplied by CMU.

Accessing Array Elements

- Get $a[i].j$
 - `sizeof(S3)` includes alignment spacers
- Element j is at offset 8 within structure
- Assembler gives offset $a+8$
 - resolved during linking

```
struct S3 {  
    short k;  
    float v;  
    short j;  
} a[10];
```



```
short get_j(int idx)  
{  
    return a[idx].j;  
}
```

```
# %eax = idx  
leal (%eax,%eax,2),%eax # 3*idx  
movswl a+8(,%eax,4),%eax
```

Supplied by CMU.

Note that even though the procedure `get_j` returns a short, its return value is converted to an int (a long word) — this is because of the convention that register `eax` (a 32-bit register) contains the return value.

Saving Space



- Put large data types first

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```

- Effect (K=4)



Supplied by CMU.

Summary

- **Arrays in C**
 - contiguous allocation of memory
 - aligned to satisfy every element's alignment requirement
 - pointer to first element
 - no bounds checking
- **Structures**
 - allocate bytes in order declared
 - pad in middle and at end to satisfy alignment

Supplied by CMU.