

CS 33

Machine Programming (6)

Most of the slides in this lecture are either from or adapted from slides provided by the authors of the textbook “Computer Systems: A Programmer’s Perspective,” 2nd Edition and are provided from the website of Carnegie-Mellon University, course 15-213, taught by Randy Bryant and David O’Hallaron in Fall 2010. These slides are indicated “Supplied by CMU” in the notes section of the slides.

Today

- **Memory Layout**
- **Buffer Overflow**
 - vulnerability
 - protection

Supplied by CMU.

Today

- **Buffer Overflow**
 - vulnerability
 - protection

Supplied by CMU.


Internet Worm and IM War

- **November, 1988**
 - Internet Worm attacks thousands of Internet hosts.
 - how did it happen?

Supplied by CMU.

Internet Worm and IM War

- November, 1988
 - Internet Worm attacks thousands of Internet hosts
 - how did it happen?
- July, 1999
 - Microsoft launches MSN Messenger (instant messaging system)
 - Messenger clients can access popular AOL Instant Messaging Service (AIM) servers



Supplied by CMU.

Internet Worm and IM War (cont.)

- **August 1999**
 - mysteriously, Messenger clients can no longer access AIM servers
 - Microsoft and AOL begin the IM war:
 - » AOL changes server to disallow Messenger clients
 - » Microsoft makes changes to clients to defeat AOL changes
 - » at least 13 such skirmishes
 - how did it happen?
- **The Internet Worm and AOL/Microsoft War were both based on *stack buffer overflow* exploits!**
 - » many library functions do not check argument sizes.
 - » allows target buffers to overflow.

Supplied by CMU.

String Library Code

- Implementation of Unix function `gets()`

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getchar();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getchar();
    }
    *p = '\0';
    return dest;
}
```

- no way to specify limit on number of characters to read
- Similar problems with other library functions
 - `strcpy`, `strcat`: copy strings of arbitrary length
 - `scanf`, `fscanf`, `sscanf`, when given `%s` conversion specification

Supplied by CMU.

Vulnerable Buffer Code

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
void call_echo() {  
    echo();  
}
```

```
unix>./bufdemo  
Type a string:1234567  
1234567
```

```
unix>./bufdemo  
Type a string:12345678  
Segmentation Fault
```

```
unix>./bufdemo  
Type a string:123456789ABC  
Segmentation Fault
```

Supplied by CMU.

Buffer Overflow Disassembly

echo:

80485c5:	55	push	%ebp
80485c6:	89 e5	mov	%esp,%ebp
80485c8:	53	push	%ebx
80485c9:	83 ec 14	sub	\$0x14,%esp
80485cc:	8d 5d f8	lea	0xffffffff8(%ebp),%ebx
80485cf:	89 1c 24	mov	%ebx,(%esp)
80485d2:	e8 9e ff ff ff	call	8048575 <gets>
80485d7:	89 1c 24	mov	%ebx,(%esp)
80485da:	e8 05 fe ff ff	call	80483e4 <puts@plt>
80485df:	83 c4 14	add	\$0x14,%esp
80485e2:	5b	pop	%ebx
80485e3:	5d	pop	%ebp
80485e4:	c3	ret	


call_echo:

80485eb:	e8 d5 ff ff ff	call	80485c5 <echo>
80485f0:	c9	leave	
80485f1:	c3	ret	

Supplied by CMU.

Buffer Overflow Stack

Before call to gets



← **%ebp**

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
echo:  
    pushl %ebp          # Save %ebp on stack  
    movl  %esp, %ebp  
    pushl %ebx          # Save %ebx  
    subl  $20, %esp     # Allocate stack space  
    leal  -8(%ebp), %ebx # Compute buf as %ebp-8  
    movl  %ebx, (%esp)   # Push buf on stack  
    call  gets          # Call gets  
    . . .
```


Supplied by CMU.

Buffer Overflow Stack Example


```

unix> gdb bufdemo
(gdb) break echo
Breakpoint 1 at 0x80485c9
(gdb) run
Breakpoint 1, 0x80485c9 in echo ()
(gdb) print /x $ebp
$1 = 0xffffd678
(gdb) print /x *(unsigned *)$ebp
$2 = 0xffffd688
(gdb) print /x *((unsigned *)$ebp + 1)
$3 = 0x80485f0
    
```

Before call to gets



Before call to gets

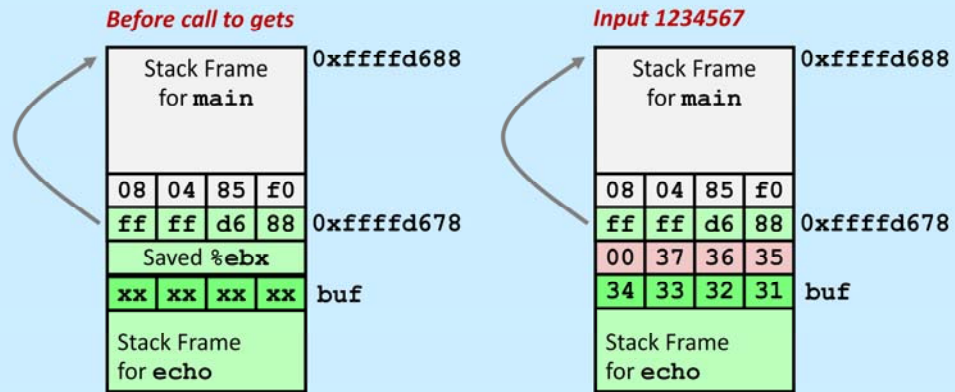


80485eb: e8 d5 ff ff ff
80485f0: c9

call 80485c5 <echo>
leave

Supplied by CMU.

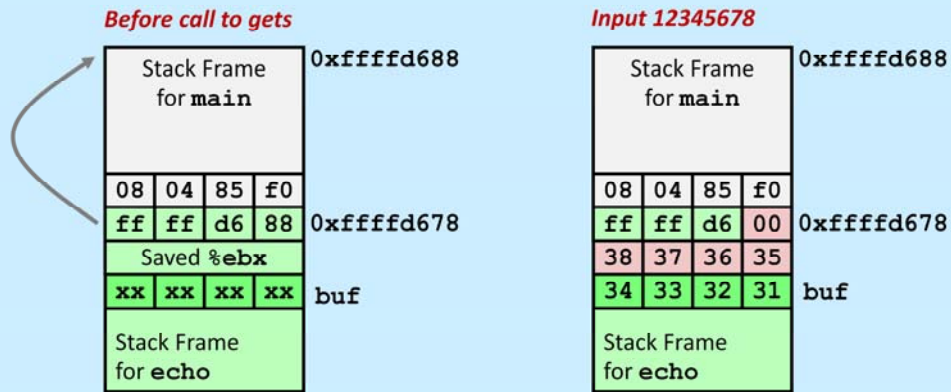
Buffer Overflow Example #1



**Overflow buf, and corrupt %ebx,
but no problem**

Supplied by CMU.

Buffer Overflow Example #2



Base pointer corrupted

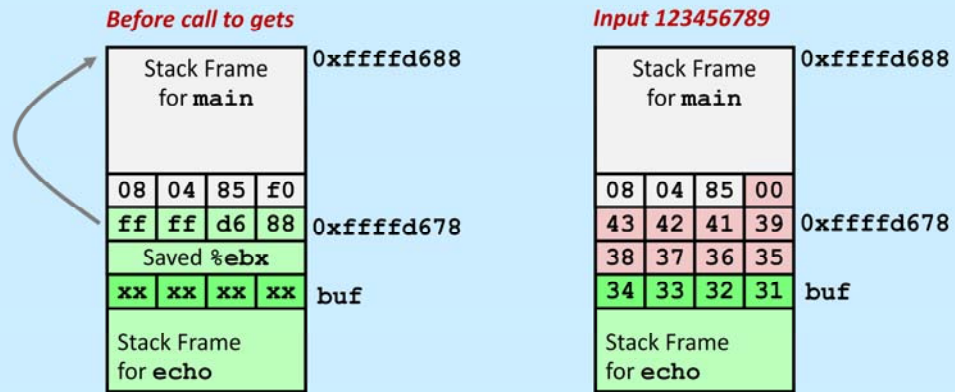
```

. . .
80485eb:  e8 d5 ff ff ff  call  80485c5 <echo>
80485f0:  c9                leave # Set %ebp to corrupted value
80485f1:  c3                ret

```

Supplied by CMU.

Buffer Overflow Example #3



Return address corrupted

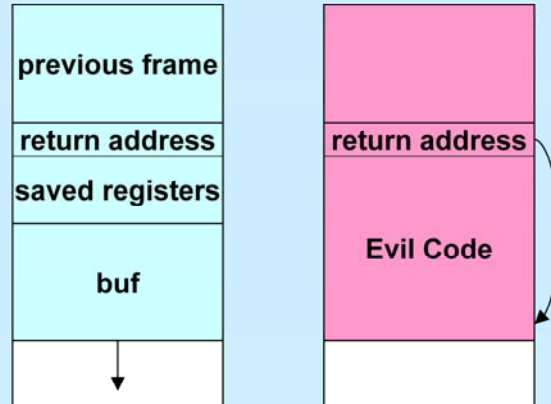
```

80485eb:  e8 d5 ff ff ff  call  80485c5 <echo>
80485f0:  c9              leave # Desired return point
    
```

Supplied by CMU.

Buffer Overflow


```
void fingerd( ) {  
    char buf[80];  
    ...  
    gets(buf);  
    ...  
}
```



Programs susceptible to buffer-overflow attacks are amazingly common and thus such attacks are probably the most common of the bug-exploitation techniques. Even drivers for network interface devices have such problems, making machines vulnerable to attacks by maliciously created packets.

Defense

```
void proc( ) {  
    char buf[80];  
    ...  
    fgets(buf, 80, stdin);  
    ...  
}
```



Malicious Use of Buffer Overflow


```
void foo() {  
    bar();  
    ...  
}
```

return
address
A

```
int bar() {  
    char buf[64];  
    gets(buf);  
    ...  
    return ...;  
}
```

data written
by `gets()`

B



- Input string contains byte representation of executable code
- Overwrite return address A with address of buffer B
- When `bar()` executes `ret`, will jump to exploit code

Supplied by CMU.

Avoiding Overflow Vulnerability

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    fgets(buf, 4, stdin);  
    puts(buf);  
}
```

- **Use library routines that limit string lengths**
 - **fgets** instead of **gets**
 - **strncpy** instead of **strcpy**
 - **don't use scanf with %s conversion specification**
 - » use **fgets** to read the string
 - » or use **%ns** where **n** is a suitable integer

Supplied by CMU.

Exploits Based on Buffer Overflows

- *Buffer overflow bugs allow remote machines to execute arbitrary code on victim machines*
- Internet worm
 - early versions of the finger server (fingerd) used `gets()` to read the argument sent by the client:
 - » `finger twd@cs.brown.edu`
 - worm attacked fingerd server by sending phony argument:
 - » `finger "exploit-code padding new-return-address"`
 - » exploit code: executed a root shell on the victim machine with a direct TCP connection to the attacker.

Supplied by CMU.

Exploits Based on Buffer Overflows

- *Buffer overflow bugs allow remote machines to execute arbitrary code on victim machines*
- **IM War**
 - AOL exploited existing buffer overflow bug in AIM clients
 - exploit code: returned 4-byte signature (the bytes at some location in the AIM client) to server
 - when Microsoft changed code to match signature, AOL changed signature location

Supplied by CMU.

Date: Wed, 11 Aug 1999 11:30:57 -0700 (PDT)
From: Phil Bucking <philbucking@yahoo.com>
Subject: AOL exploiting buffer overrun bug in their own software!
To: rms@pharlap.com

Mr. Smith,

I am writing you because I have discovered something that I think you might find interesting because you are an Internet security expert with experience in this area. I have also tried to contact AOL but received no response.

I am a developer who has been working on a revolutionary new instant messaging client that should be released later this year.

...

It appears that the AIM client has a buffer overrun bug. By itself this might not be the end of the world, as MS surely has had its share. But AOL is now *exploiting their own buffer overrun bug* to help in its efforts to block MS Instant Messenger.

....

Since you have significant credibility with the press I hope that you can use this information to help inform people that behind AOL's friendly exterior they are nefariously compromising peoples' security.

Sincerely,
Phil Bucking
Founder, Bucking Consulting
philbucking@yahoo.com

***It was later determined that this
email originated from within
Microsoft!***

Supplied by CMU.

Code-Red Exploit Code

- Attacked Microsoft IIS servers
- Starts 100 threads running
- Spreads self
 - generate random IP addresses & send attack string
 - between 1st & 19th of month
- Attack www.whitehouse.gov
 - send 98,304 packets; sleep for 4-1/2 hours: repeat
 - » denial-of-service attack
 - between 21st & 27th of month
- Deface server's home page
 - after waiting 2 hours



Supplied by CMU.

A detailed description of the exploit can be found at <http://www.eeye.com/Resources/Security-Center/Research/Security-Advisories/AL20010717>.

System-Level Protections

- **Randomized stack offsets**
 - at start of program, allocate random amount of space on stack
 - makes it difficult for hacker to predict beginning of inserted code
- **Nonexecutable code segments**
 - in traditional x86, can mark region of memory as either “read-only” or “writeable”
 - » can execute anything readable
 - modern hardware requires explicit “execute” permission

```
unix> gdb bufdemo
(gdb) break echo

(gdb) run
(gdb) print /x $ebp
$1 = 0xffffc638

(gdb) run
(gdb) print /x $ebp
$2 = 0xffffbb08

(gdb) run
(gdb) print /x $ebp
$3 = 0xffffc6a8
```

Stack Canaries

- **Idea**
 - place special value (“canary”) on stack just beyond buffer
 - check for corruption before exiting function
- **gcc implementation**
 - `-fstack-protector`
 - `-fstack-protector-all`

```
unix>./bufdemo-protected
Type a string:1234
1234
```

```
unix>./bufdemo-protected
Type a string:12345
*** stack smashing detected ***
```

Supplied by CMU.

Protected Buffer Disassembly echo:

```
804864d: 55                push    %ebp
804864e: 89 e5            mov     %esp,%ebp
8048650: 53              push    %ebx
8048651: 83 ec 14        sub     $0x14,%esp
8048654: 65 a1 14 00 00 00 mov     %gs:0x14,%eax
804865a: 89 45 f8        mov     %eax,0xffffffff8(%ebp)
804865d: 31 c0          xor     %eax,%eax
804865f: 8d 5d f4        lea     0xffffffff4(%ebp),%ebx
8048662: 89 1c 24        mov     %ebx,(%esp)
8048665: e8 77 ff ff ff  call    80485e1 <gets>
804866a: 89 1c 24        mov     %ebx,(%esp)
804866d: e8 ca fd ff ff  call    804843c <puts@plt>
8048672: 8b 45 f8        mov     0xffffffff8(%ebp),%eax
8048675: 65 33 05 14 00 00 00 xor     %gs:0x14,%eax
804867c: 74 05          je      8048683 <echo+0x36>
804867e: e8 a9 fd ff ff  call    804842c <FAIL>
8048683: 83 c4 14        add     $0x14,%esp
8048686: 5b              pop     %ebx
8048687: 5d              pop     %ebp
8048688: c3              ret
```

Supplied by CMU.

The operand “%gs:0x14” requires some explanation, as it uses features we haven’t previously discussed. gs is one of a few “segment registers,” which refer to other areas of memory. They are generally not used, being a relic of the early days of the x86 architecture before virtual-memory support was added. You can think of it as an area where global variables (accessible from anywhere) may be stored and made read-only. It’s used here to store the “canary” values.

Setting Up Canary

Before call to gets




```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
echo:  
    . . .  
    movl    %gs:20, %eax    # Get canary  
    movl    %eax, -8(%ebp)  # Put on stack  
    xorl    %eax, %eax     # Erase canary  
    . . .
```

Supplied by CMU.

Checking Canary

Before call to gets



← %ebp

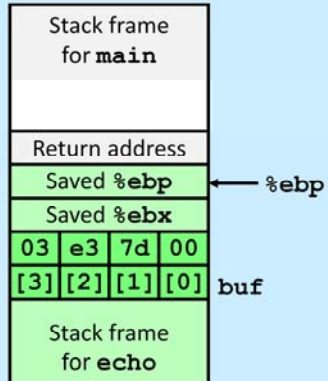
```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

```
echo:
    . . .
    movl    -8(%ebp), %eax    # Retrieve from stack
    xorl    %gs:20, %eax     # Compare with Canary
    je      .L24             # Same: skip ahead
    call    __stack_chk_fail # ERROR
.L24:
    . . .
```

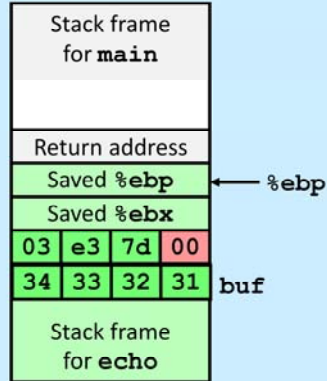
Supplied by CMU.

Canary Example

Before call to gets



Input 1234



```
(gdb) break echo
(gdb) run
(gdb) stepi 3
(gdb) print /x *((unsigned *) $ebp - 2)
$1 = 0x3e37d00
```

Supplied by CMU.

Today

- **Buffer Overflow**
 - vulnerability
 - protection

Supplied by CMU.