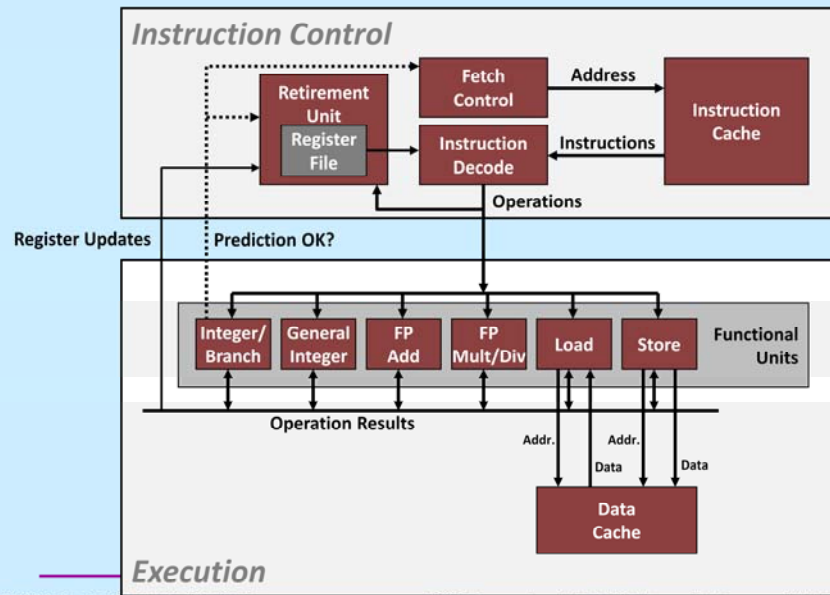


# CS 33

## Architecture and Optimization (3)

Some of the slides in this lecture are either from or adapted from slides provided by the authors of the textbook “Computer Systems: A Programmer’s Perspective,” 2<sup>nd</sup> Edition and are provided from the website of Carnegie-Mellon University, course 15-213, taught by Randy Bryant and David O’Hallaron in Fall 2010. These slides are indicated “Supplied by CMU” in the notes section of the slides.

# Modern CPU Design



Supplied by CMU.

## Latency of Loads

```
typedef struct ELE {  
    struct ELE *next;  
    int data;  
} list_ele, *list_ptr;
```

```
int list_len(list_ptr ls) {  
    int len = 0;  
    while (ls) {  
        len++;  
        ls = ls->next;  
    }  
    return len;  
}
```

```
# len in %eax, ls in %rdi
```


```
.L11:                # loop:  
    addl    $1, %eax    # incr len  
    movq    (%rdi), %rdi # ls - ls->next  
    testq   %rdi, %rdi  # test ls  
    jne     .L11        # if != 0  
                    # go to loop
```

• 4 CPE

This example is from the textbook (Figure 5.31). The point is that loads (fetching data from memory) have a latency of 4 cycles.


## Data Flow Over Multiple Iterations

- Why doesn't load fit into the critical path?




This is Figure 5.15 of Bryant and O'Hallaron.

## Data-Flow Graphs of Inner Loop



These are Figures 5.14 a and b of Bryant and O'Hallaron.

## Parallel Execution



Here we look in detail at the timing requirements of the example of the previous slides. In the current slide, the length of a box indicates how much latency the operation requires: *mul* and *load* require 4 cycles, *add* requires 1. The first operations that can execute are the *load* and the *add*. Once the *load* completes, the *mul* operation may start, since it's waiting on the value loaded. However, note that the next instance of *load* may start as soon as *add* has updated *%rdx*. Since loads are pipelined, it thus may start execution one cycle after the first *load* started. The next instance of *add* may start as soon as all users of *%rdx* (in this case, the *load* and *cmp* instructions) are done with *%rdx*. Thus it certainly may execute anytime after the second *load* completes.

The second *mul* may not start until after the first *mul* has produced its result (in *%xmm0*) and after *load* has loaded the multiplier. But since the second *load* produced the multiplier for the second *mul* shortly after the first *mul* began, the second *mul* must wait only until the first *mul* completes.

## Clearing an Array ...

```
#define ITERS 100000000
int main() {
    volatile int dest[100];
    int iter;
    for (iter=0; iter<ITERS; iter++) {
        long i;
        for (i=0; i<100; i++)
            dest[i] = 0;
    }
}
```

This is adapted from Figure 5.32 from the textbook.

## Clearing an Array ... Unwound

```
#define ITERS 100000000
int main() {
    volatile int dest[100];
    int iter;
    for (iter=0; iter<ITERS; iter++) {
        long i;
        for (i=0; i<96; i+=4) {
            dest[i] = 0;
            dest[i+1] = 0;
            dest[i+2] = 0;
            dest[i+3] = 0;
        }
    }
}
```

This is adapted from Figure 5.32 from the textbook.



## Store/Load Interaction

```
void write_read(int *src, int *dest, int n) {  
    int cnt = n;  
    int val = 0;  
  
    while(cnt-->0) {  
        *dest = val;  
        val = (*src)+1;  
    }  
}
```

This code is from the textbook.

## Store/Load Interaction

Example A: `write_read(&a[0], &a[1], 3)`

|     | Initial  | Iter. 1 | Iter. 2  | Iter. 3  |
|-----|----------|---------|----------|----------|
| cnt | 3        | 2       | 1        | 0        |
| a   | -10   17 | -10   0 | -10   -9 | -10   -9 |
| val | 0        | -9      | -9       | -9       |

• CPE 2.0


Example B: `write_read(&a[0], &a[0], 3)`

|     | Initial  | Iter. 1 | Iter. 2 | Iter. 3 |
|-----|----------|---------|---------|---------|
| cnt | 3        | 2       | 1       | 0       |
| a   | -10   17 | 0   17  | 1   17  | 2   17  |
| val | 0        | 1       | 2       | 3       |

• CPE 6.0

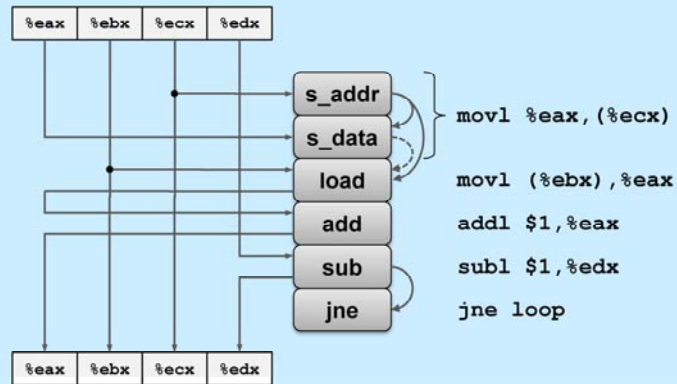
This is Figure 5.33 from the textbook.

## Some Details of Load and Store



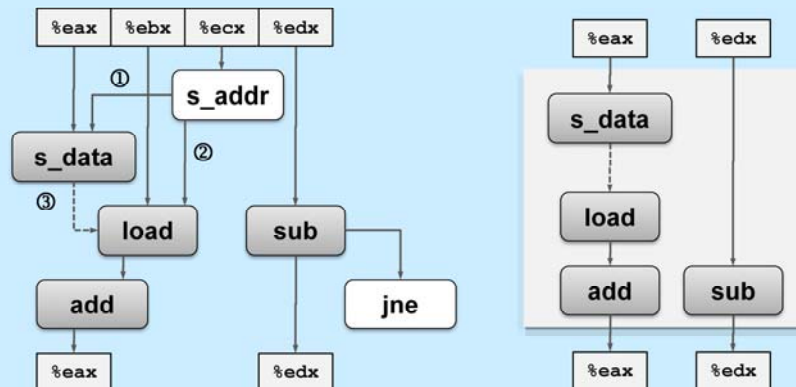
This is Figure 5.34 from the textbook.

## Inner-Loop Data Flow of Write\_Read



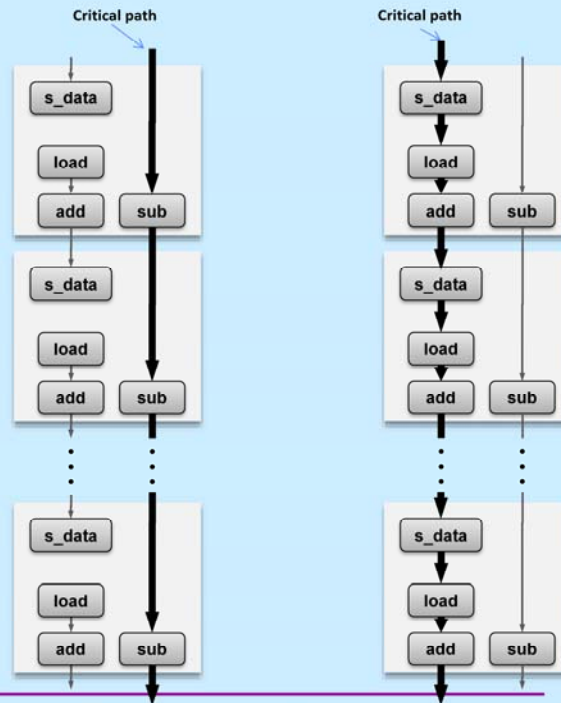
This is Figure 5.35 from the textbook.

## Inner-Loop Data Flow of Write\_Read



This is Figure 5.36 from the textbook.

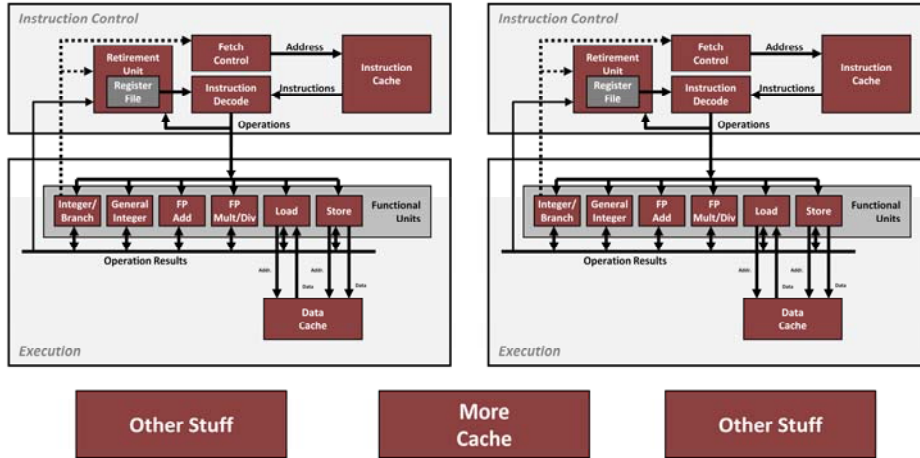
## Data Flow



This is adapted from Figure 5.37 from the textbook.

# Multiple Cores

## Chip



# Hyper Threading

