

CS 33

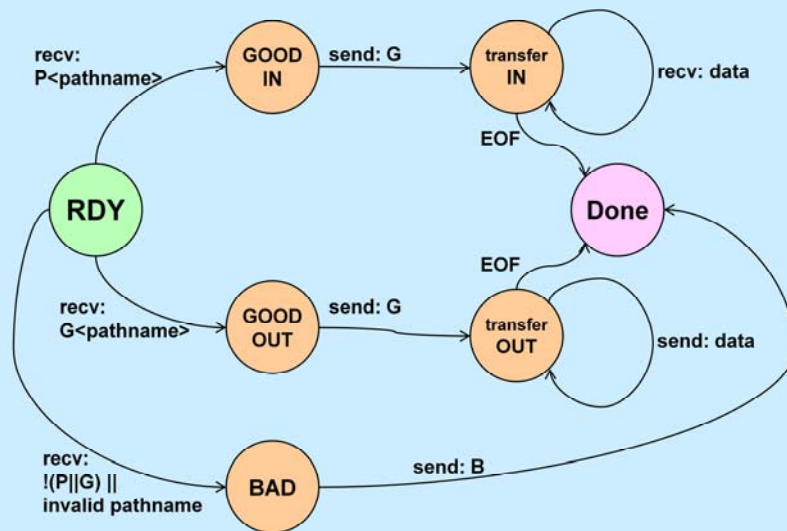
More Network Programming

The source code used in this lecture is on the course web page.

A Really Simple Protocol

- **Transfer a file**
 - layered on top of TCP
 - » reliable
 - » indicates if connection is closed
- **To send a file**
 - P<null-terminated pathname><contents of file>
- **To retrieve a file**
 - G<null-terminated pathname><contents of file>

Server State Machine



Keeping Track of State

```
typedef struct client {
    int fd;        // file descriptor of local file being transferred
    int size;      // size of out-going data in buffer
    char buf[BSIZE];
    enum state {RDY, BAD, GOOD, TRANSFER} state;
    /*
       states:
       RDY: ready to receive client's command (P or G)
       BAD: client's command was bad, sending B response + error msg
       GOOD: client's command was good, sending G response
       TRANSFER: transferring data
    */
    enum dir {IN, OUT} dir;
    /*
       IN: client has issued P command
       OUT: client has issued G command
    */
} client_t;
```

Main Server Loop

```
while(1) {
    select(maxfd, &trd, &twr, 0, 0);
    if (FD_ISSET(lsock, &trd)) {
        // a new connection
        new_client(lsock);
    }
    for (i=lsock+1; i<maxfd; i++) {
        if (FD_ISSET(i, &trd)) {
            // ready to read
            read_event(i);
        }
        if (FD_ISSET(i, &twr)) {
            // ready to write
            write_event(i);
        }
    }
    trd = rd; twr = wr;
}
```

Note that *trd*, *twr*, *rd* and *wr* are allof type *fd_set*. *rd* and *wr* are initialized so that *rd* contains just the file descriptor for the listening socket and *wr* is empty. *trd* and *twr* are copied from *rd* and *wr* respectively before the loop is entered.

New Client

```
// Accept a new connection on listening socket  
// fd. Return the connected file descriptor
```

```
int new_client(int fd) {  
    int cfd = accept(fd, 0, 0);  
    clients[cfd].state = RDY;  
    FD_SET(cfd, &rd);  
    return cfd;  
}
```

Read Event (1)

```
// File descriptor fd is ready to be read. Read it, then handle
// the input
void read_event(int fd) {
    client_t *c = &clients[fd];
    int ret = read(fd, c->buf, BSIZE);
    switch (c->state) {
    case RDY:
        if (c->buf[0] == 'G') {
            // GET request (to fetch a file)
            c->dir = OUT;
            if ((c->fd = open(&c->buf[1], O_RDONLY)) == -1) {
                // open failed; send negative response and error message
                c->state = BAD;
                c->buf[0] = 'B';
                strncpy(&c->buf[1], strerror(errno)+1, BSIZE-2);
                c->buf[BSIZE-1] = 0;
                c->size = strlen(c->buf)+1;
            }
        }
    }
```

Read Event (2)

```
    else {  
        // open succeeded; send positive response  
        c->state = GOOD;  
        c->size = 1;  
        c->buf[0] = 'G';  
    }  
    // prepare to send response to client  
    FD_SET(fd, &wr);  
    FD_CLR(fd, &rd);  
    break;  
}
```


Read Event (3)

```
if (c->buf[0] == 'P') {
    // PUT request (to create a file)
    c->dir = IN;
    if ((c->fd = open(&c->buf[1],
        O_RDWR|O_CREAT|O_TRUNC, 0666)) == -1) {
        // open failed; send negative response and error message
        ...
    } else {
        // open succeeded; send positive response
        ...
    }
    // prepare to send response to client
    FD_SET(fd, &wr);
    FD_CLR(fd, &rd);
    break;
}
```

Read Event (4)

```
case TRANSFER:
    // should be in midst of receiving file contents from client
    if (ret == 0) {
        // eof: all done
        close(c->fd);
        close(fd);
        FD_CLR(fd, &rd);
        break;
    }
    if (write(c->fd, c->buf, ret) == -1) {
        // write to file failed: terminate connection to client
        ...
        break;
    }
    // continue to read more data from client
    break;
}
```

Write Event (1)

```
// File descriptor fd is ready to be written to. Write to it, then,  
// depending on current state, prepare for the next action.  
void write_event(int fd) {  
    client_t *c = &clients[fd];  
    int ret = write(fd, c->buf, c->size);  
    if (ret == -1) {  
        // couldn't write to client; terminate connection  
        close(c->fd);  
        close(fd);  
        FD_CLR(fd, &wr);  
        c->fd = -1;  
        perror("write to client");  
        return;  
    }  
    switch (c->state) {
```

Write Event (2)

```
case BAD:
    // finished sending error message; now terminate client connection
    close(c->fd);
    close(fd);
    FD_CLR(fd, &wr);
    c->fd = -1;
break;
```

Write Event (3)

```
case GOOD:
    if (c->dir == IN) {
        // finished response to PUT request
        c->state = TRANSFER;
        FD_SET(fd, &rd);
        FD_CLR(fd, &wr);
        break;
    }
    // otherwise finished response to GET request, so proceed
```


Write Event (4)

```
case TRANSFER:
    // should be in midst of transferring file contents to client
    if ((c->size = read(c->fd, c->buf, BSIZE)) == -1) {
        ...
        break;
    } else if (c->size == 0) {
        // no more file to transfer; terminate client connection
        close(c->fd);
        close(fd);
        FD_CLR(fd, &wr);
        c->fd = -1;
        break;
    }
    // continue to write more data to client
    break;
}
```

Problems


- **Works fine as long as protocol is followed correctly**
 - can client (malicious or incompetent) cause server to misbehave?
- **How does server limit file access?**

chroot (before)



The `chroot` system call allows a process to change the root of its file-system hierarchy to a directory lower down in the tree.

chroot (after)



After a successful invocation of *chroot*, paths starting with “/” are followed starting at the new root.

Secure chroot?

- **Implementation**
 - “..” = “.” at process's root
 - » can't *cd* to parent
- **Secure?**
 - leakproof?

No ...

```
chdir("/");  
pfd = open(".", O_RDONLY);  
mkdir("Houdini", 0700);  
chroot("Houdini");  
fchdir(pfd);  
for (i=0; i<100; i++)  
    chdir("../");  
chroot(".");
```

The problem is that current open files are not closed, even if they are outside of the subtree to which access is restricted.

The *fchdir* system call changes the current directory to be the directory referred to by the file descriptor passed as an argument.

chroot

- **Would work fine for our simple file transfer protocol**
 - actually is used in *tftp* (trivial file transfer protocol)
 - however, requires superuser powers (so as to help avoid problems of previous slide)

